

01. Data Exploration and Feature Engineering

Project Overview

Objective: Detect fraudulent Medicare providers using claims and beneficiary data.

Datasets:

- `Train_Beneficiarydata.csv` : Patient demographics and chronic conditions
- `Train_Inpatientdata.csv` : Hospital admission claims
- `Train_Outpatientdata.csv` : Outpatient visit claims
- `Train_Labels.csv` : Provider-level fraud labels (Yes/No)

Key Challenge: Class imbalance (~10% fraud rate) and multi-table data requiring aggregation.

```
In [1]: # Import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime
import warnings
warnings.filterwarnings('ignore')

# Visualization Settings
sns.set_style('whitegrid')
plt.rcParams['figure.figsize'] = (12, 6)
plt.rcParams['font.size'] = 10

print("Libraries imported successfully!")
```

Libraries imported successfully!

1. Load Data

```
In [2]: DATA_DIR = '../data/'

# Load datasets
train_bene = pd.read_csv(DATA_DIR + 'Train_Beneficiarydata.csv')
train_inpatient = pd.read_csv(DATA_DIR + 'Train_Inpatientdata.csv')
train_outpatient = pd.read_csv(DATA_DIR + 'Train_Outpatientdata.csv')
train_labels = pd.read_csv(DATA_DIR + 'Train_Labels.csv')

print(f"Beneficiary Data Shape: {train_bene.shape}")
print(f"Inpatient Data Shape: {train_inpatient.shape}")
print(f"Outpatient Data Shape: {train_outpatient.shape}")
```

```

print(f"Labels Data Shape: {train_labels.shape}")

print("\nFirst few rows of Labels:")
train_labels.head()

```

Beneficiary Data Shape: (138556, 25)
 Inpatient Data Shape: (40474, 30)
 Outpatient Data Shape: (517737, 27)
 Labels Data Shape: (5410, 2)

First few rows of Labels:

	Provider	PotentialFraud
0	PRV51001	No
1	PRV51003	Yes
2	PRV51004	No
3	PRV51005	Yes
4	PRV51007	No

2. Understanding Data Relationships

Key Identifiers

1. BenID (Beneficiary ID):

- Links patients to their claims
- One beneficiary can have multiple claims (both inpatient and outpatient)
- Found in: `Train_Beneficiarydata`, `Train_Inpatientdata`, `Train_Outpatientdata`

2. Provider ID:

- Links claims to fraud labels
- One provider can serve multiple beneficiaries and file multiple claims
- Found in: `Train_Inpatientdata`, `Train_Outpatientdata`, `Train_Labels`

3. ClaimID:

- Unique identifier for each claim
- Found in: `Train_Inpatientdata`, `Train_Outpatientdata`

Unit of Analysis: PROVIDER

Why Provider-Level?

- Fraud labels are assigned at the **Provider** level, not claim or patient level
- A fraudulent provider may file many claims, some legitimate and some fraudulent

- We must aggregate all claim-level and patient-level information to characterize each provider's behavior

Data Flow:

```

Beneficiary Data (BeneID)
    ↓ (merge on BeneID)
Claims Data (BeneID, Provider, ClaimID)
    ↓ (aggregate by Provider)
Provider-Level Features (Provider)
    ↓ (merge on Provider)
Final Dataset with Labels (Provider, Features, PotentialFraud)

```

3. Exploratory Data Analysis (EDA)

3.1 Target Variable Distribution

```
In [3]: # Class distribution
fraud_counts = train_labels['PotentialFraud'].value_counts()
fraud_pct = train_labels['PotentialFraud'].value_counts(normalize=True) * 100

print("Fraud Distribution:")
print(fraud_counts)
print("\nPercentages:")
print(fraud_pct)

# Visualization
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Count plot
sns.countplot(data=train_labels, x='PotentialFraud', palette='Set2', ax=axes[0])
axes[0].set_title('Fraud Label Distribution (Count)', fontsize=14, fontweight='bold')
axes[0].set_ylabel('Count')
for p in axes[0].patches:
    axes[0].annotate(f'{int(p.get_height())}', 
                    (p.get_x() + p.get_width() / 2., p.get_height()),
                    ha='center', va='bottom')

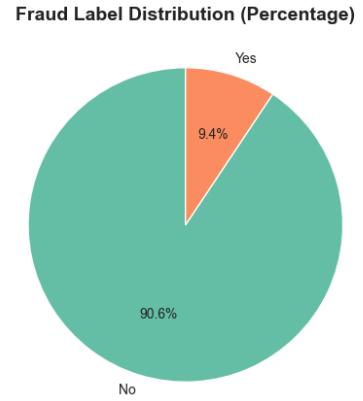
# Pie chart
axes[1].pie(fraud_counts, labels=fraud_counts.index, autopct='%1.1f%%',
            colors=['#66c2a5', '#fc8d62'], startangle=90)
axes[1].set_title('Fraud Label Distribution (Percentage)', fontsize=14, fontweight='bold')

plt.tight_layout()
plt.show()

print(f"\n⚠ CLASS IMBALANCE DETECTED: {fraud_pct['Yes']:.2f}% fraud cases")
```

```
Fraud Distribution:  
PotentialFraud  
No      4904  
Yes     506  
Name: count, dtype: int64
```

```
Percentages:  
PotentialFraud  
No      90.64695  
Yes     9.35305  
Name: proportion, dtype: float64
```



⚠ CLASS IMBALANCE DETECTED: 9.35% fraud cases

3.2 Missing Value Analysis

```
In [4]: def analyze_missing(df, name):  
    """Analyze and visualize missing values"""  
    missing = df.isnull().sum()  
    missing_pct = (missing / len(df)) * 100  
    missing_df = pd.DataFrame({  
        'Column': missing.index,  
        'Missing_Count': missing.values,  
        'Missing_Percentage': missing_pct.values  
    })  
    missing_df = missing_df[missing_df['Missing_Count'] > 0].sort_values('Missing_P  
  
    if len(missing_df) > 0:  
        print(f"\n{'='*60}")  
        print(f"Missing Values in {name}")  
        print(f"{'='*60}")  
        print(missing_df.to_string(index=False))  
  
        # Plot top 15 columns with missing values  
        if len(missing_df) > 0:  
            plt.figure(figsize=(12, 6))  
            top_missing = missing_df.head(15)  
            sns.barplot(data=top_missing, y='Column', x='Missing_Percentage', palette  
            plt.title(f'Top Missing Values in {name}', fontsize=14, fontweight='bold')  
            plt.xlabel('Missing Percentage (%)')  
            plt.tight_layout()  
            plt.show()
```

```

    else:
        print(f"\n✓ No missing values in {name}")

# Analyze each dataset
analyze_missing(train_bene, "Beneficiary Data")
analyze_missing(train_inpatient, "Inpatient Data")
analyze_missing(train_outpatient, "Outpatient Data")
analyze_missing(train_labels, "Labels Data")

```


=====

Missing Values in Beneficiary Data

=====

Column	Missing_Count	Missing_Percentage
DOD	137135	98.974422

Top Missing Values in Beneficiary Data

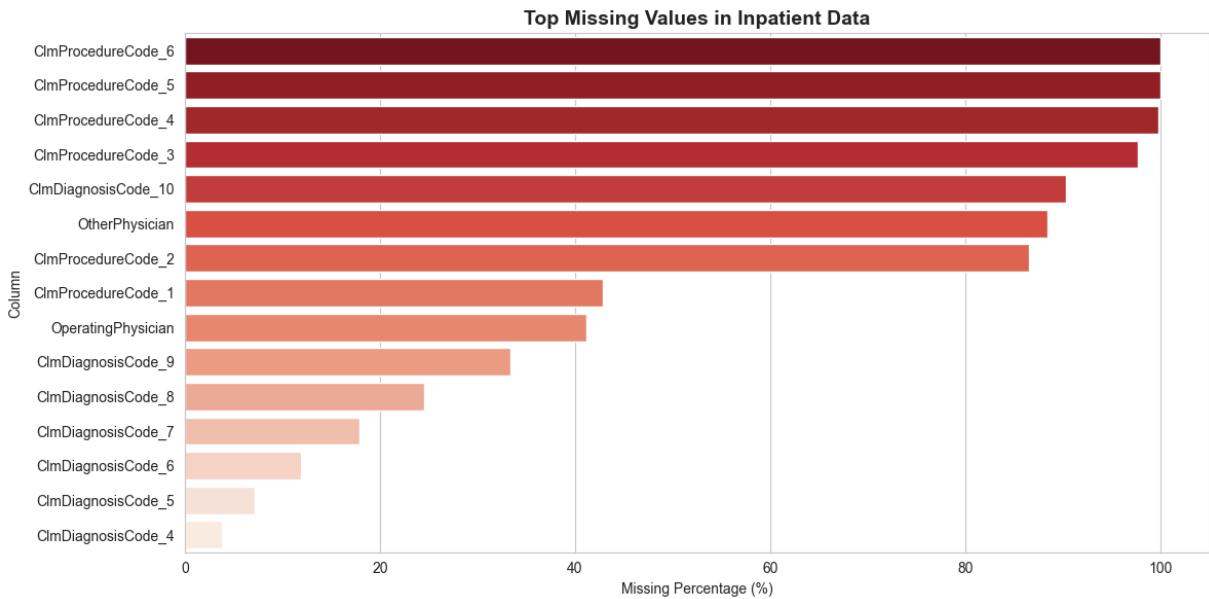


=====

Missing Values in Inpatient Data

=====

Column	Missing_Count	Missing_Percentage
ClmProcedureCode_6	40474	100.000000
ClmProcedureCode_5	40465	99.977764
ClmProcedureCode_4	40358	99.713396
ClmProcedureCode_3	39509	97.615753
ClmDiagnosisCode_10	36547	90.297475
OtherPhysician	35784	88.412314
ClmProcedureCode_2	35020	86.524683
ClmProcedureCode_1	17326	42.807728
OperatingPhysician	16644	41.122696
ClmDiagnosisCode_9	13497	33.347334
ClmDiagnosisCode_8	9942	24.563918
ClmDiagnosisCode_7	7258	17.932500
ClmDiagnosisCode_6	4838	11.953353
ClmDiagnosisCode_5	2894	7.150269
ClmDiagnosisCode_4	1534	3.790087
DeductibleAmtPaid	899	2.221179
ClmDiagnosisCode_3	676	1.670208
ClmDiagnosisCode_2	226	0.558383
AttendingPhysician	112	0.276721

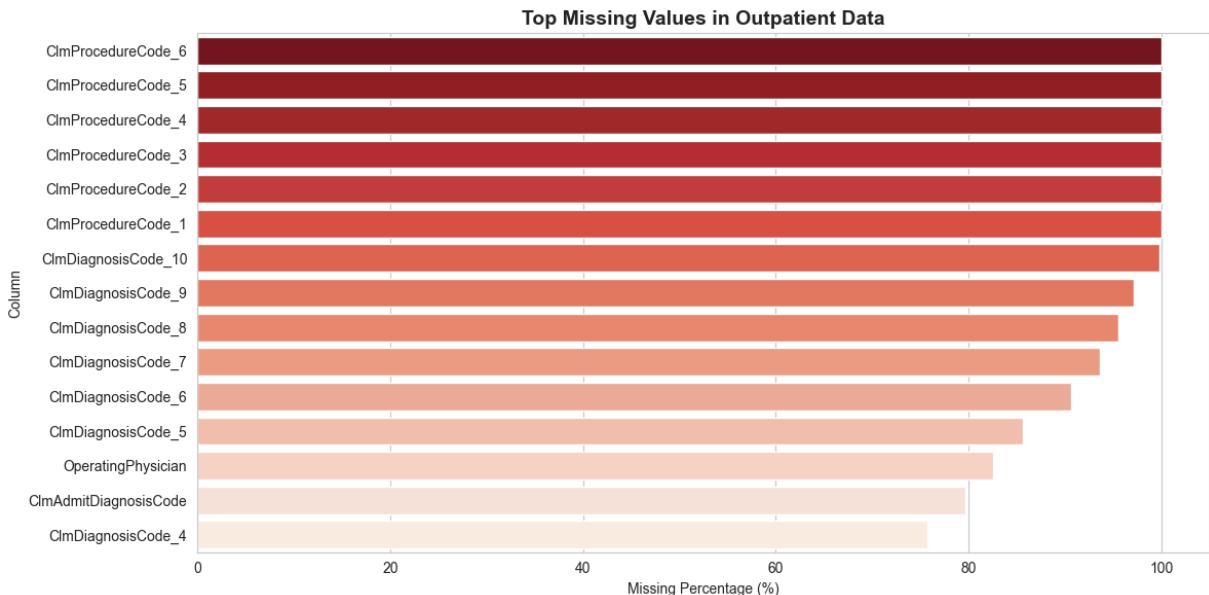


=====

Missing Values in Outpatient Data

=====

Column	Missing_Count	Missing_Percentage
ClmProcedureCode_6	517737	100.000000
ClmProcedureCode_5	517737	100.000000
ClmProcedureCode_4	517735	99.999614
ClmProcedureCode_3	517733	99.999227
ClmProcedureCode_2	517701	99.993047
ClmProcedureCode_1	517575	99.968710
ClmDiagnosisCode_10	516654	99.790820
ClmDiagnosisCode_9	502899	97.134066
ClmDiagnosisCode_8	494825	95.574587
ClmDiagnosisCode_7	484776	93.633640
ClmDiagnosisCode_6	468981	90.582864
ClmDiagnosisCode_5	443393	85.640586
OperatingPhysician	427120	82.497484
ClmAdmitDiagnosisCode	412312	79.637345
ClmDiagnosisCode_4	392141	75.741351
OtherPhysician	322691	62.327205
ClmDiagnosisCode_3	314480	60.741264
ClmDiagnosisCode_2	195380	37.737307
ClmDiagnosisCode_1	10453	2.018979
AttendingPhysician	1396	0.269635



✓ No missing values in Labels Data

Observations:

- Diagnosis and procedure codes have high missingness (expected - not all claims use all code fields)
- Physician fields have some missingness
- `DeductibleAmtPaid` has missing values (likely when fully covered by insurance)

3.3 Data Preprocessing

```
In [5]: # Convert date columns
date_cols = ['ClaimStartDt', 'ClaimEndDt', 'AdmissionDt', 'DischargeDt']

for col in date_cols:
    if col in train_inpatient.columns:
        train_inpatient[col] = pd.to_datetime(train_inpatient[col], errors='coerce')
    if col in train_outpatient.columns:
        train_outpatient[col] = pd.to_datetime(train_outpatient[col], errors='coerce')

train_bene['DOB'] = pd.to_datetime(train_bene['DOB'], errors='coerce')
train_bene['DOD'] = pd.to_datetime(train_bene['DOD'], errors='coerce')

# Calculate claim duration
train_inpatient['ClaimDuration'] = (train_inpatient['ClaimEndDt'] - train_inpatient['ClaimStartDt'])
train_outpatient['ClaimDuration'] = (train_outpatient['ClaimEndDt'] - train_outpatient['ClaimStartDt'])

# Calculate length of stay for inpatient
if 'AdmissionDt' in train_inpatient.columns and 'DischargeDt' in train_inpatient.columns:
    train_inpatient['LengthOfStay'] = (train_inpatient['DischargeDt'] - train_inpatient['AdmissionDt']).dt.days

# Calculate age (approximate at 2009)
train_bene['Age'] = 2009 - train_bene['DOB'].dt.year

print("✓ Date preprocessing completed")
```

```

print(f"Inpatient claim duration range: {train_inpatient['ClaimDuration'].min()} to {train_inpatient['ClaimDuration'].max()}")
print(f"Outpatient claim duration range: {train_outpatient['ClaimDuration'].min()} to {train_outpatient['ClaimDuration'].max()}")

```

✓ Date preprocessing completed
Inpatient claim duration range: 1 to 37 days
Outpatient claim duration range: 1 to 24 days

3.4 Merge Data for Analysis

```

In [6]: # Merge beneficiary info into claims
inpatient_full = pd.merge(train_inpatient, train_bene, on='BeneID', how='left')
outpatient_full = pd.merge(train_outpatient, train_bene, on='BeneID', how='left')

# Merge Labels for EDA
inpatient_eda = pd.merge(inpatient_full, train_labels, on='Provider', how='left')
outpatient_eda = pd.merge(outpatient_full, train_labels, on='Provider', how='left')

print(f"Inpatient (with beneficiary & labels): {inpatient_eda.shape}")
print(f"Outpatient (with beneficiary & labels): {outpatient_eda.shape}")

```

Inpatient (with beneficiary & labels): (40474, 58)
Outpatient (with beneficiary & labels): (517737, 54)

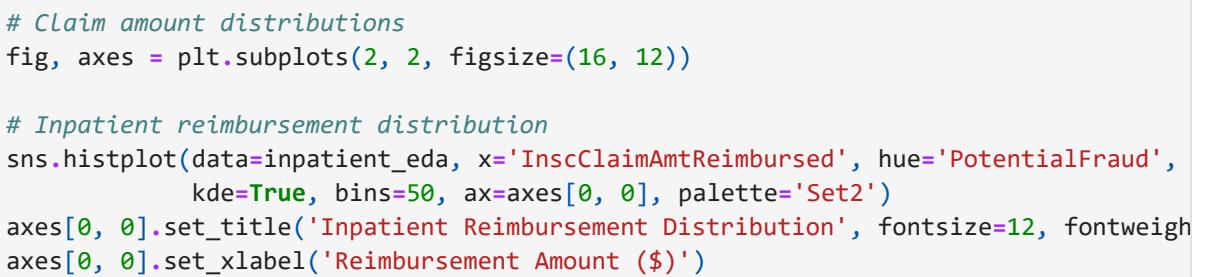
3.5 Financial Features Analysis

```

In [7]: # Claim amount distributions
fig, axes = plt.subplots(2, 2, figsize=(16, 12))

# Inpatient reimbursement distribution
sns.histplot(data=inpatient_eda, x='InscClaimAmtReimbursed', hue='PotentialFraud',
              kde=True, bins=50, ax=axes[0, 0], palette='Set2')
axes[0, 0].set_title('Inpatient Reimbursement Distribution', fontsize=12, fontweight='bold')
axes[0, 0].set_xlabel('Reimbursement Amount ($)')

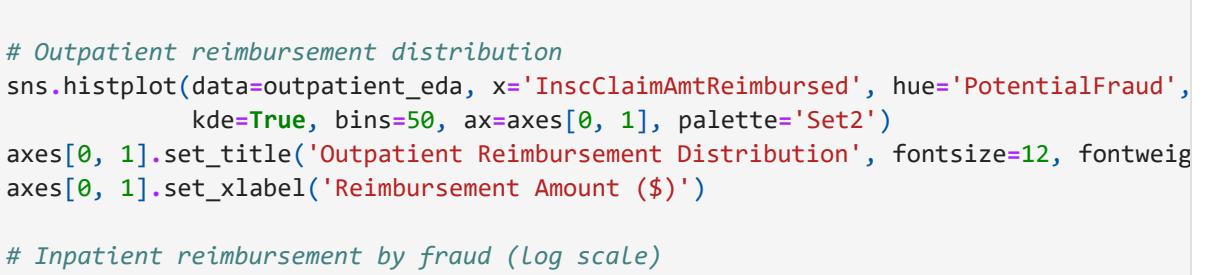
```



```

# Outpatient reimbursement distribution
sns.histplot(data=outpatient_eda, x='InscClaimAmtReimbursed', hue='PotentialFraud',
              kde=True, bins=50, ax=axes[0, 1], palette='Set2')
axes[0, 1].set_title('Outpatient Reimbursement Distribution', fontsize=12, fontweight='bold')
axes[0, 1].set_xlabel('Reimbursement Amount ($)')

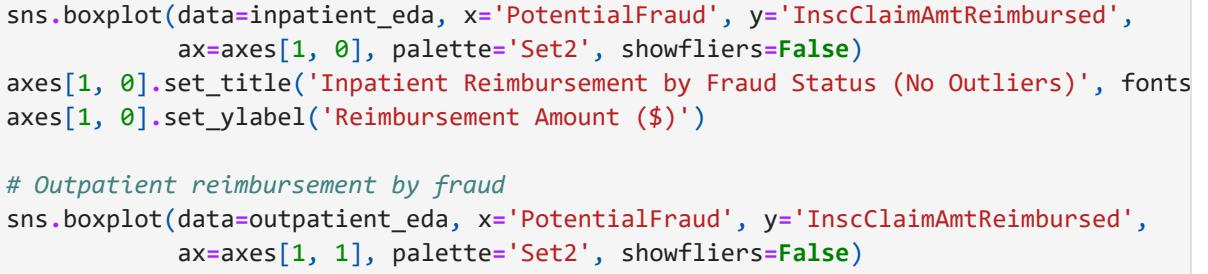
```



```

# Inpatient reimbursement by fraud (log scale)
sns.boxplot(data=inpatient_eda, x='PotentialFraud', y='InscClaimAmtReimbursed',
            ax=axes[1, 0], palette='Set2', showfliers=False)
axes[1, 0].set_title('Inpatient Reimbursement by Fraud Status (No Outliers)', fontsize=12, fontweight='bold')
axes[1, 0].set_ylabel('Reimbursement Amount ($)')

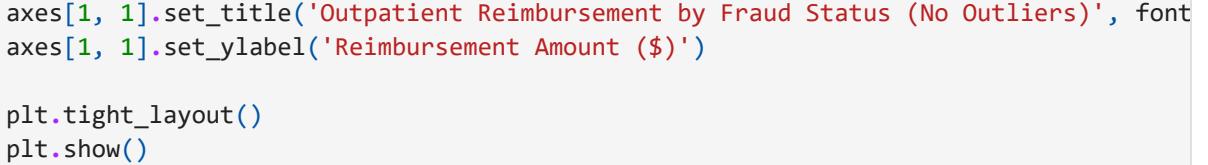
```



```

# Outpatient reimbursement by fraud
sns.boxplot(data=outpatient_eda, x='PotentialFraud', y='InscClaimAmtReimbursed',
            ax=axes[1, 1], palette='Set2', showfliers=False)
axes[1, 1].set_title('Outpatient Reimbursement by Fraud Status (No Outliers)', fontsize=12, fontweight='bold')
axes[1, 1].set_ylabel('Reimbursement Amount ($)')

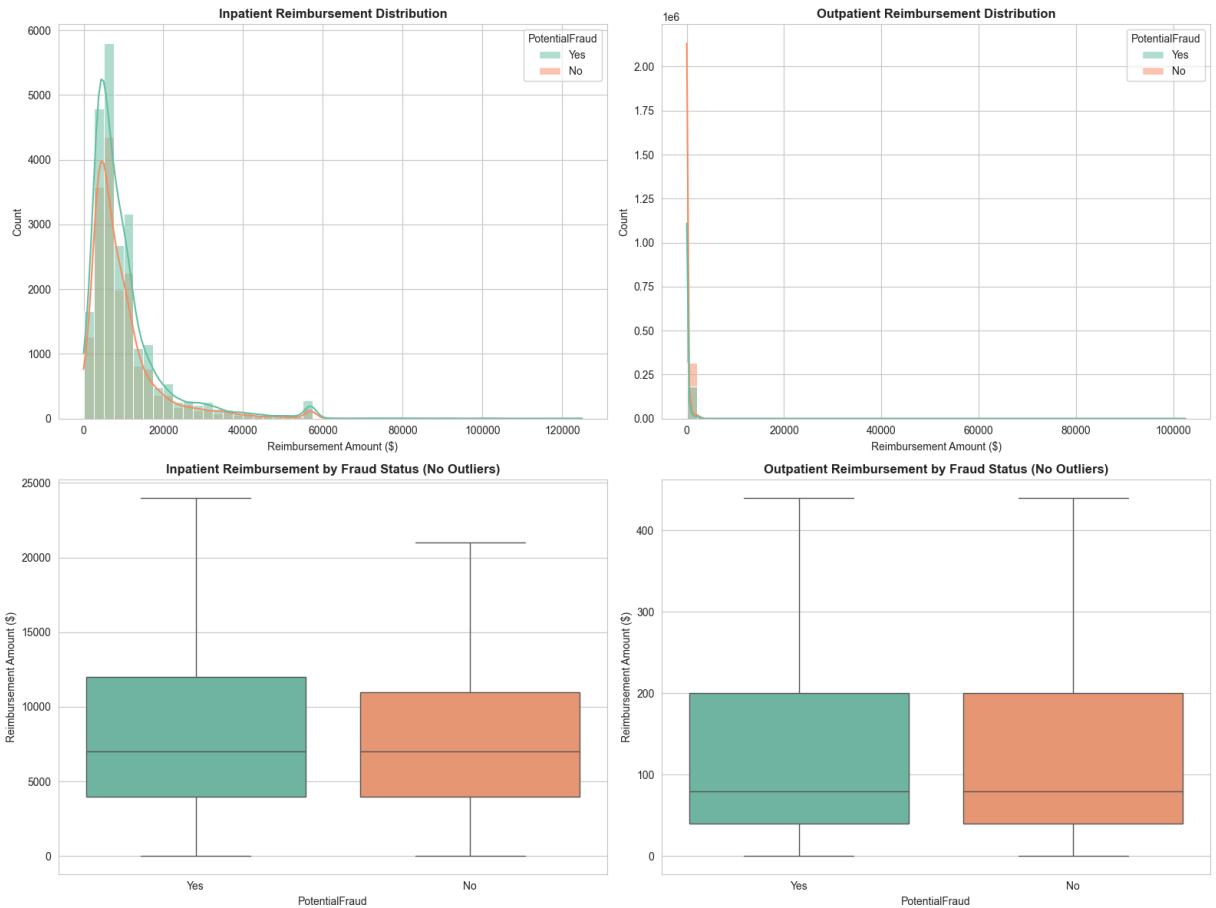
```



```

plt.tight_layout()
plt.show()

```



3.6 Claim Duration Analysis

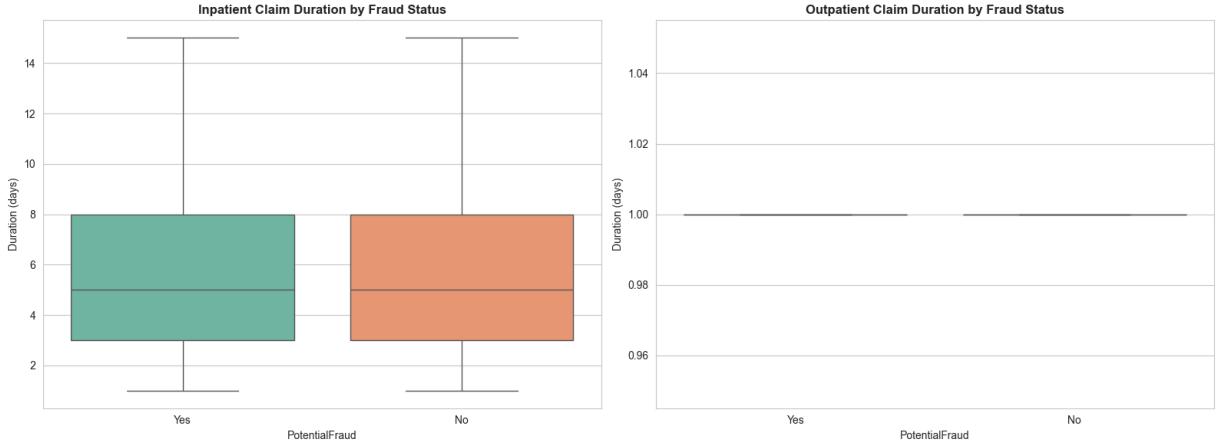
```
In [8]: fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# Inpatient claim duration
sns.boxplot(data=inpatient_eda, x='PotentialFraud', y='ClaimDuration',
            ax=axes[0], palette='Set2', showfliers=False)
axes[0].set_title('Inpatient Claim Duration by Fraud Status', fontsize=12, fontweight='bold')
axes[0].set_ylabel('Duration (days)')

# Outpatient claim duration
sns.boxplot(data=outpatient_eda, x='PotentialFraud', y='ClaimDuration',
            ax=axes[1], palette='Set2', showfliers=False)
axes[1].set_title('Outpatient Claim Duration by Fraud Status', fontsize=12, fontweight='bold')
axes[1].set_ylabel('Duration (days)')

plt.tight_layout()
plt.show()

# Statistical comparison
print("\nClaim Duration Statistics:")
print("\nInpatient:")
print(inpatient_eda.groupby('PotentialFraud')[['ClaimDuration']].describe())
print("\nOutpatient:")
print(outpatient_eda.groupby('PotentialFraud')[['ClaimDuration']].describe())
```



Claim Duration Statistics:

Inpatient:

	count	mean	std	min	25%	50%	75%	max
PotentialFraud								
No	17072.0	6.553831	5.329792	1.0	3.0	5.0	8.0	36.0

Outpatient:

	count	mean	std	min	25%	50%	75%	max
PotentialFraud								
No	328343.0	2.413022	4.695344	1.0	1.0	1.0	1.0	22.0

3.7 Outlier Detection

```
In [9]: # Identify outliers using IQR method
def detect_outliers(df, column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    outliers = df[(df[column] < lower_bound) | (df[column] > upper_bound)]
    return len(outliers), (len(outliers) / len(df)) * 100

print("Outlier Analysis:")
print("\nInpatient Reimbursement:")
count, pct = detect_outliers(inpatient_eda, 'InscClaimAmtReimbursed')
print(f"  Outliers: {count} ({pct:.2f}%)")

print("\nOutpatient Reimbursement:")
count, pct = detect_outliers(outpatient_eda, 'InscClaimAmtReimbursed')
print(f"  Outliers: {count} ({pct:.2f}%)")

print("\n⚠️ High outlier percentage suggests diverse provider types and potential +")
```

Outlier Analysis:

Inpatient Reimbursement:

Outliers: 2966 (7.33%)

Outpatient Reimbursement:

Outliers: 78375 (15.14%)

⚠️ High outlier percentage suggests diverse provider types and potential fraud patterns

4. Feature Engineering: Provider-Level Aggregation

Strategy

We aggregate claim-level and beneficiary-level data to create comprehensive provider profiles:

1. **Financial Features:** Total/average reimbursements, deductibles
2. **Volume Features:** Claim counts, beneficiary counts, claim ratios
3. **Medical Complexity:** Diagnosis/procedure counts, chronic conditions
4. **Temporal Features:** Average claim durations, length of stay
5. **Demographic Features:** Average patient age

```
In [10]: def engineer_provider_features(claims_df, prefix):  
    """  
        Aggregate claim-level data to provider level  
  
        Parameters:  
        - claims_df: DataFrame with claims and merged beneficiary data  
        - prefix: 'Inpatient' or 'Outpatient'  
  
        Returns:  
        - DataFrame with provider-level features  
    """  
  
    # Chronic condition columns (convert 2=No to 0, keep 1=Yes)  
    chronic_cols = [col for col in claims_df.columns if 'ChronicCond' in col]  
    for col in chronic_cols:  
        claims_df[col] = claims_df[col].replace(2, 0)  
  
    # Count diagnosis and procedure codes  
    diag_cols = [col for col in claims_df.columns if 'ClmDiagnosisCode' in col]  
    proc_cols = [col for col in claims_df.columns if 'ClmProcedureCode' in col]  
  
    claims_df['NumDiagnoses'] = claims_df[diag_cols].notnull().sum(axis=1)  
    claims_df['NumProcedures'] = claims_df[proc_cols].notnull().sum(axis=1)  
  
    # Aggregation functions  
    agg_dict = {  
        # Volume  
        'ClaimID': 'count',
```

```

'BeneID': 'nunique',

# Financial
'InscClaimAmtReimbursed': ['sum', 'mean', 'std', 'max'],
'DeductibleAmtPaid': ['sum', 'mean'],

# Temporal
'ClaimDuration': ['mean', 'max'],

# Demographics
'Age': 'mean',

# Medical complexity
'NumDiagnoses': ['mean', 'max'],
'NumProcedures': ['mean', 'max'],

# Physicians
'AttendingPhysician': 'nunique',
'OperatingPhysician': 'nunique',
'OtherPhysician': 'nunique'
}

# Add Length of stay for inpatient
if 'LengthOfStay' in claims_df.columns:
    agg_dict['LengthOfStay'] = ['mean', 'max']

# Add chronic conditions (mean = percentage of patients with condition)
for col in chronic_cols:
    agg_dict[col] = 'mean'

# Group by provider
provider_features = claims_df.groupby('Provider').agg(agg_dict)

# Flatten column names
provider_features.columns = [
    f"{prefix}_{col[0]}_{col[1]}" if isinstance(col, tuple) else f"{prefix}_{col}"
    for col in provider_features.columns
]

return provider_features

print("Aggregating Inpatient features...")
inpatient_features = engineer_provider_features(inpatient_full, 'Inpatient')
print(f"  Shape: {inpatient_features.shape}")

print("\nAggregating Outpatient features...")
outpatient_features = engineer_provider_features(outpatient_full, 'Outpatient')
print(f"  Shape: {outpatient_features.shape}")

```

Aggregating Inpatient features...

Shape: (2092, 31)

Aggregating Outpatient features...

Shape: (5012, 29)

4.1 Combine Features and Create Derived Metrics

```
In [11]: # Merge inpatient and outpatient features
provider_df = pd.merge(inpatient_features, outpatient_features,
                      on='Provider', how='outer')

# Fill NaN with 0 (providers with only inpatient or only outpatient)
provider_df = provider_df.fillna(0)

print(f"Combined provider features shape: {provider_df.shape}")

# Create derived features
print("\nCreating derived features...")

# Total claims
provider_df['Total_Claims'] = (provider_df['Inpatient_ClaimID_count'] +
                                 provider_df['Outpatient_ClaimID_count'])

# Total reimbursement
provider_df['Total_Reimbursement'] = (provider_df['Inpatient_InscClaimAmtReimbursed'] +
                                         provider_df['Outpatient_InscClaimAmtReimbursed'])

# Average reimbursement per claim
provider_df['Avg_Reimbursement_Per_Claim'] = provider_df['Total_Reimbursement'] / (provider_df['Total_Claims'])

# Inpatient to outpatient ratio
provider_df['Inpatient_Outpatient_Ratio'] = (provider_df['Inpatient_ClaimID_count'] /
                                              provider_df['Outpatient_ClaimID_count'])

# Total unique beneficiaries
provider_df['Total_Unique_Beneficiaries'] = (provider_df['Inpatient_BeneID_nunique'] +
                                               provider_df['Outpatient_BeneID_nunique'])

# Claims per beneficiary
provider_df['Claims_Per_Beneficiary'] = provider_df['Total_Claims'] / (provider_df['Total_Unique_Beneficiaries'])

# Average chronic conditions (sum of percentages)
chronic_in_cols = [c for c in provider_df.columns if 'ChronicCond' in c and 'Inpatient' in c]
chronic_out_cols = [c for c in provider_df.columns if 'ChronicCond' in c and 'Outpatient' in c]

if chronic_in_cols:
    provider_df['Avg_Chronic_Conditions_Inpatient'] = provider_df[chronic_in_cols].sum(1) / len(chronic_in_cols)
if chronic_out_cols:
    provider_df['Avg_Chronic_Conditions_Outpatient'] = provider_df[chronic_out_cols].sum(1) / len(chronic_out_cols)

print(f"\n✓ Feature engineering complete. Total features: {provider_df.shape[1]}")
```

Combined provider features shape: (5410, 60)

Creating derived features...

✓ Feature engineering complete. Total features: 68

4.2 Merge with Labels

```
In [12]: # Merge with fraud Labels
provider_df_final = pd.merge(provider_df, train_labels, on='Provider', how='inner')

# Encode target variable
provider_df_final['PotentialFraud'] = provider_df_final['PotentialFraud'].map({'Yes': 1, 'No': 0})

print(f"Final dataset shape: {provider_df_final.shape}")
print(f"\nTarget distribution:")
print(provider_df_final['PotentialFraud'].value_counts())

provider_df_final.head()
```

Final dataset shape: (5410, 70)

Target distribution:

PotentialFraud	count
0	4904
1	506

Name: count, dtype: int64

```
Out[12]:   Provider Inpatient_ClaimID_count Inpatient_BenelD_nunique Inpatient_InscClaimAmtReimbursement
0 PRV51001           5.0                  5.0
1 PRV51003          62.0                 53.0
2 PRV51004           0.0                  0.0
3 PRV51005           0.0                  0.0
4 PRV51007           3.0                  3.0
```

5 rows × 70 columns

4.3 Provider-Level Descriptive Statistics

```
In [13]: # Key statistics by fraud status
key_features = ['Total_Claims', 'Total_Reimbursement', 'Total_Unique_Beneficiaries',
                'Avg_Reimbursement_Per_Claim', 'Inpatient_Outpatient_Ratio']

print("Provider-Level Statistics by Fraud Status:")
print("*80")
for feature in key_features:
    print(f"\n{feature}:")
    print(provider_df_final.groupby('PotentialFraud')[feature].describe())
```

Provider-Level Statistics by Fraud Status:

Total_Claims:

PotentialFraud	count	mean	std	min	25%	50%	75%	\
0	4904.0	70.435359	128.942510	1.0	9.0	27.0	72.0	
1	506.0	420.545455	722.734485	1.0	62.0	155.5	432.0	

max

PotentialFraud

0	1245.0
1	8240.0

Total_Reimbursement:

PotentialFraud	count	mean	std	min	25%	\
0	4904.0	53193.723491	102342.349409	0.0	3797.5	
1	506.0	584350.039526	644668.507561	200.0	172947.5	

50% 75% max

PotentialFraud

0	15055.0	57422.5	1311040.0
1	373450.0	759740.0	5996050.0

Total_Underlying_Beneficiaries:

	count	mean	std	min	25%	50%	75%	\
PotentialFraud								
0	4904.0	49.385808	82.375614	1.0	8.0	22.0	54.00	
1	506.0	247.527668	349.867285	1.0	52.0	120.5	252.75	
	max							
PotentialFraud								
0		807.0						
1		2857.0						
Avg_Reimbursement_Per_Claim:								\
PotentialFraud	count	mean	std	min	25%			
0	4904.0	1523.777469	3375.550421	0.000000	220.498226			
1	506.0	3842.793675	3812.438253	73.333211	857.060180			
	50%	75%		max				
PotentialFraud								
0		332.192040	1024.733692	56999.430006				
1		2576.479936	5759.701690	22333.258889				
Inpatient_Outpatient_Ratio:								\
PotentialFraud	count	mean	std	min	25%	50%	75%	\
0	4904.0	0.655777	2.955304	0.0	0.000000	0.000000	0.085791	
1	506.0	3.013342	9.851957	0.0	0.064804	0.276053	1.020045	
	max							
PotentialFraud								
0		50.0						
1		119.0						

4.4 Correlation Analysis

```
In [14]: # Select numeric columns for correlation
numeric_cols = provider_df_final.select_dtypes(include=[np.number]).columns.tolist()


# Correlation with target
correlations = provider_df_final[numeric_cols].corr()['PotentialFraud'].sort_values
print("Top 15 Features Correlated with Fraud:")
print(correlations.head(15))

# Correlation heatmap (top features)
top_features = correlations.abs().sort_values(ascending=False).head(20).index.tolist
plt.figure(figsize=(14, 12))
sns.heatmap(provider_df_final[top_features].corr(), annot=True, fmt=' .2f',
            cmap='coolwarm', center=0, square=True, linewidths=1)
plt.title('Correlation Heatmap: Top 20 Features', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.show()
```

Top 15 Features Correlated with Fraud:

PotentialFraud	1.000000
Total_Reimbursement	0.575558
Inpatient_ClaimDuration_max	0.542879
Inpatient_LengthOfStay_max	0.542662
Inpatient_InscClaimAmtReimbursed_sum	0.532795
Inpatient_DeductibleAmtPaid_sum	0.525454
Inpatient_ClaimID_count	0.525393
Inpatient_BeneID_nunique	0.522256
Inpatient_InscClaimAmtReimbursed_max	0.504854
Inpatient_NumProcedures_max	0.444011
Total_Unique_Beneficiaries	0.399033
Total_Claims	0.374197
Inpatient_InscClaimAmtReimbursed_std	0.364039
Inpatient_NumDiagnoses_max	0.342265
Outpatient_BeneID_nunique	0.340550

Name: PotentialFraud, dtype: float64



5. Save Final Dataset

```
In [15]: # Save to CSV
output_path = DATA_DIR + 'provider_features.csv'
provider_df_final.to_csv(output_path, index=False)

print(f"✓ Final dataset saved to: {output_path}")
print(f"  Shape: {provider_df_final.shape}")
print(f"  Features: {provider_df_final.shape[1] - 2}")
print(f"  Providers: {provider_df_final.shape[0]}")
print(f"  Fraud cases: {provider_df_final['PotentialFraud'].sum()}")
```

✓ Final dataset saved to: ../data/provider_features.csv
Shape: (5410, 70)
Features: 68
Providers: 5410
Fraud cases: 506

Summary

Key Findings from EDA:

1. **Class Imbalance:** ~10% fraud rate requires special handling
2. **Missing Data:** High missingness in diagnosis/procedure codes (expected)
3. **Financial Patterns:** Fraudulent providers show different reimbursement distributions
4. **Outliers:** Significant outliers in financial features suggest diverse fraud patterns

Feature Engineering Accomplishments:

- Created **comprehensive provider-level features** from claim and beneficiary data
- Engineered **financial, volume, medical complexity, and temporal features**
- Generated **derived metrics** (ratios, averages, totals)
- Final dataset ready for modeling

Next Steps:

→ Proceed to **Notebook 02: Modeling**